

```

from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

import os
import time
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
!pip install tensorflow_addons
import tensorflow_addons as tfa
from keras.models import Sequential, Model
from keras import backend as K
from keras.layers import LeakyReLU
from keras.activations import selu
from keras.layers import Multiply, Add
from keras.optimizers import Adam
from keras.layers import Dense, Conv2D, Conv2DTranspose, Input, Flatten, BatchNormalization, Lambda, Reshape, Activation

Collecting tensorflow_addons
  Downloading tensorflow_addons-0.22.0-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (612 kB)
    612.3/612.3 kB 8.7 MB/s eta 0:00:00
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from tensorflow_addons) (23.2)
Collecting typeguard<3.0.0,>=2.7 (from tensorflow_addons)
  Downloading typeguard-2.13.3-py3-none-any.whl (17 kB)
Installing collected packages: typeguard, tensorflow_addons
Successfully installed tensorflow_addons-0.22.0 typeguard-2.13.3
/usr/local/lib/python3.10/dist-packages/tensorflow_addons/utils/tfa_eol_msg.py:23: UserWarning:

TensorFlow Addons (TFA) has ended development and introduction of new features.
TFA has entered a minimal maintenance and release mode until a planned end of life in May 2024.
Please modify downstream libraries to take dependencies from other repositories in our TensorFlow community (e.g. Keras, Ke

For more information see: https://github.com/tensorflow/addons/issues/2807

warnings.warn(

train_dir = "/content/drive/MyDrive/Dogs/dogs_data/train"
train_images = [os.path.join(train_dir, image) for image in os.listdir(train_dir)]
print(len(train_images))
val_dir = "/content/drive/MyDrive/Dogs/dogs_data/val"
val_images = [os.path.join(val_dir, image) for image in os.listdir(val_dir)]
print(len(val_images))
test_dir = "/content/drive/MyDrive/Dogs/dogs_data/test"
test_images = [os.path.join(test_dir, image) for image in os.listdir(test_dir)]
print(len(test_images))

810
101
101

# preprocess
image_size = 64

def preprocess(image):
    # Read and decode the image
    image = tf.io.read_file(image)
    image = tf.io.decode_jpeg(image)

    # if augment:
    #     # Data Augmentation
    #     # Randomly flip horizontally
    #     image = tf.image.random_flip_left_right(image)

    #     angle = tf.random.uniform([], -20, 20, dtype=tf.float32)
    #     image = tfa.image.rotate(image, angle)

    #     # Randomly adjust brightness
    #     image = tf.image.random_brightness(image, max_delta=0.2)
    #     # You can change the max_delta value to control the range of brightness adjustment.

    #     # Randomly zoom in or out (scale between 0.9 and 1.1)
    #     scale = tf.random.uniform([], 0.9, 1.1, dtype=tf.float32)
    #     new_height = tf.cast(image_size * scale, dtype=tf.int32)

```

```
# new_width = tf.cast(image_size * scale, dtype=tf.int32)
# image = tf.image.resize(image, (new_height, new_width))

# Randomly shear the image
# shear = tf.random.uniform([], -0.2, 0.2, dtype=tf.float32)
# image = tf.keras.preprocessing.image.apply_affine_transform(image, shear=shear)

# Cast and resize the image
image = tf.cast(image, tf.float32)
image = tf.image.resize(image, (image_size, image_size))
image = image / 255.0
image = tf.reshape(image, shape=(image_size, image_size, 3))
return image

batch_size = 128

training_dataset = tf.data.Dataset.from_tensor_slices((train_images))
# training_dataset = training_dataset.map(lambda x: preprocess(x, augment=False))
training_dataset = training_dataset.map(preprocess)

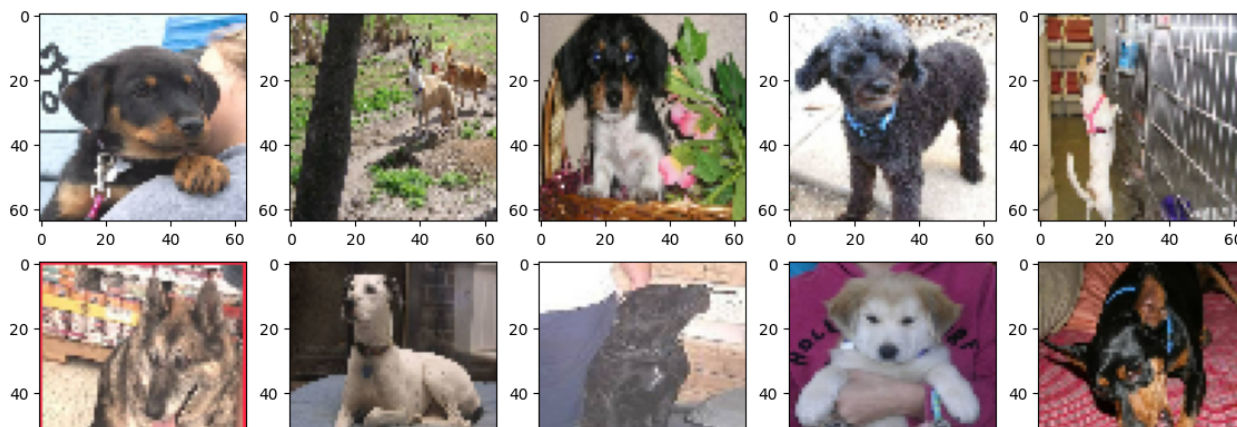
training_dataset = training_dataset.shuffle(1000).batch(batch_size)

val_dataset = tf.data.Dataset.from_tensor_slices((val_images))
val_dataset = val_dataset.map(preprocess)
val_dataset = val_dataset.shuffle(1000).batch(batch_size)

test_dataset = tf.data.Dataset.from_tensor_slices((test_images))
test_dataset = test_dataset.map(preprocess)
test_dataset = test_dataset.shuffle(1000).batch(batch_size)

# visualize some of them
fig, axes = plt.subplots(5,5, figsize = (14,14))
sample = training_dataset.unbatch().take(25)
sample = [image for image in sample]

idx = 0
for row in range(5):
    for column in range(5):
        axes[row, column].imshow(sample[idx])
        idx+=1
```



```

latent_dim = 512
encoder_input = Input(shape = (64,64,3))

x = Conv2D(32, kernel_size=5, activation = LeakyReLU(0.02), strides = 1, padding = 'same')(encoder_input)
x = BatchNormalization()(x)

filter_size = [64,128,256,512]
for i in filter_size:
    x = Conv2D(i, kernel_size=5, activation = LeakyReLU(0.02), strides = 2, padding = 'same')(x)
    x = BatchNormalization()(x)

x = Flatten()(x)
x = Dense(1024, activation = selu)(x)
encoder_output = BatchNormalization()(x)

# sampling layer
mu = Dense(latent_dim)(encoder_output)
log_var = Dense(latent_dim)(encoder_output)

epsilon = K.random_normal(shape = (tf.shape(mu)[0], tf.shape(mu)[1]))
sigma = tf.exp(0.5 * log_var)

z_eps = Multiply()(sigma, epsilon)
z = Add()(mu, z_eps)

encoder = Model(encoder_input, outputs = [mu, log_var, z], name = 'encoder')
# encoder.summary()

# build the decoder
decoder = Sequential()
decoder.add(Dense(1024, activation = selu, input_shape = (latent_dim, )))
decoder.add(BatchNormalization())

decoder.add(Dense(8192, activation = selu))
decoder.add(Reshape((4,4,512)))

decoder.add(Conv2DTranspose(256, (5,5), activation = LeakyReLU(0.02), strides = 2, padding = 'same'))
decoder.add(BatchNormalization())

decoder.add(Conv2DTranspose(128, (5,5), activation = LeakyReLU(0.02), strides = 2, padding = 'same'))
decoder.add(BatchNormalization())

decoder.add(Conv2DTranspose(64, (5,5), activation = LeakyReLU(0.02), strides = 2, padding = 'same'))
decoder.add(BatchNormalization())

decoder.add(Conv2DTranspose(32, (5,5), activation = LeakyReLU(0.02), strides = 2, padding = 'same'))
decoder.add(BatchNormalization())

decoder.add(Conv2DTranspose(3, (5,5), activation = "sigmoid", strides = 1, padding = 'same'))
decoder.add(BatchNormalization())

# decoder.summary()

# vae loss = reconstruction loss + KL div

def reconstruction_loss(y, y_pred):
    return tf.reduce_mean(tf.square(y - y_pred))

```

```

def kl_loss(mu, log_var):
    loss = -0.5 * tf.reduce_mean(1 + log_var - tf.square(mu) - tf.exp(log_var))
    return loss

def vae_loss(y_true, y_pred, mu, log_var):
    return reconstruction_loss(y_true, y_pred) + (1 / (64*64)) * kl_loss(mu, log_var)

mu, log_var, z = encoder(encoder_input)
reconstructed = decoder(z)
model = Model(encoder_input, reconstructed, name="vae")
loss = kl_loss(mu, log_var)
model.add_loss(loss)
# model.summary()

checkpoint_path = '/content/drive/MyDrive/model/200e_vae_checkpoint.h5'
parameter_path = '/content/drive/MyDrive/model/200e_vae_checkpoint.txt'

model_checkpoint_callback = tf.keras.callbacks.ModelCheckpoint(
    filepath=checkpoint_path,
    save_weights_only=True,
    save_best_only=False,
    save_freq='epoch')

train_losses = []
val_losses = []
current_epoch = 0

if os.path.exists(checkpoint_path) and os.path.exists(parameter_path):
    model = tf.keras.models.load_model(checkpoint_path)

    with open(parameter_path, 'r') as file:
        lines = file.readlines()
        current_epoch = int(lines[0].split(':')[1])
        train_losses = [float(loss) for loss in lines[1].split(':')[1].strip('\n').split(',')]
        val_losses = [float(loss) for loss in lines[2].split(':')[1].strip('\n').split(',')]

    WARNING:tensorflow:No training configuration found in the save file, so the model was *not* compiled. Compile it manually.

random_vector = tf.random.normal(shape = (25, latent_dim,))

optimizer = Adam(0.001, 0.5)
epochs = 100

for epoch in range(current_epoch+1, epochs + 1):
    s_time = time.time()
    t_loss = 0.0
    for step, training_batch in enumerate(training_dataset):
        with tf.GradientTape() as tape:
            reconstructed = model(training_batch)
            y_true = tf.reshape(training_batch, shape = [-1])
            y_pred = tf.reshape(reconstructed, shape = [-1])

            mse_loss = reconstruction_loss(y_true, y_pred)

            kl = sum(model.losses)

            train_loss = 0.01 * kl + mse_loss
            t_loss += train_loss.numpy()

            grads = tape.gradient(train_loss, model.trainable_variables)
            optimizer.apply_gradients(zip(grads, model.trainable_variables))

    train_losses.append(t_loss)
    v_loss = 0.0
    for step, validation_batch in enumerate(val_dataset):
        reconstructed = model(validation_batch)
        y_true = tf.reshape(validation_batch, shape=[-1])
        y_pred = tf.reshape(reconstructed, shape=[-1])

        mse_loss = reconstruction_loss(y_true, y_pred)

        kl = sum(model.losses)

        validation_loss = 0.01 * kl + mse_loss

```

```

v_loss += validation_loss.numpy()

v_loss = (v_loss / len(val_dataset)) * len(training_dataset)
val_losses.append(v_loss)
current_epoch = epoch

model.save(checkpoint_path)

with open(parameter_path, 'w') as file:
    file.write(f'Epoch: {current_epoch}\n')
    file.write(f'Train_losses: {train_losses}\n')
    file.write(f'Val_losses: {val_losses}\n')

print("Epoch: %s - Train_loss: %.4f - Validation_loss: %.4f - Time: %.4f" % (epoch, t_loss, v_loss, time.time() - s_time))

plt.figure()
plt.plot(train_losses, label='Train Loss')
plt.plot(val_losses, label='Validation Loss')
plt.legend()
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training and Validation Loss')
plt.savefig("/content/drive/MyDrive/model/decay_lr_loss")
plt.show()

```



```

def display_original_and_reconstructed_images(original, reconstructed):
    plt.figure(figsize=(15, 5))
    num_images = 5

    for i in range(num_images):
        plt.subplot(2, num_images, i + 1)
        plt.imshow(original[i].numpy(), cmap='gray')
        plt.title("Original")
        plt.axis('off')

        plt.subplot(2, num_images, i + 1 + num_images)
        plt.imshow(reconstructed[i].numpy(), cmap='gray')
        plt.title("Reconstructed")
        plt.axis('off')

    plt.savefig('/content/drive/MyDrive/model/decay_lr_test')
    plt.show()

original_images = []
reconstructed_images = []

for test_batch in training_dataset:
    reconstructed_batch = model(test_batch)
    original_images.append(test_batch)
    reconstructed_images.append(reconstructed_batch)

```

```
# Combine the original and reconstructed images for visualization
original_images = tf.concat(original_images, axis=0)
reconstructed_images = tf.concat(reconstructed_images, axis=0)

# Display the original and reconstructed images
display_original_and_reconstructed_images(original_images, reconstructed_images)
```

